

Carving Shapes with Ruled Surfaces for Rough Machining

Paper ID: 245

Abstract

We propose a novel method for constructing a small set of ruled surfaces, each strictly outside the input shape, for rough machining. To generate such ruled surfaces, our algorithm consists of two phases: (1) generate a small set of covers for the input shape and (2) fit each cover patch using ruled surfaces without colliding with the input shape. Since the normal changes fastest along one of the principal curvature directions and the ruled surfaces cannot be deformed along their rulings, making the rulings and the principal curvature direction where the normal changes fastest orthogonal is more likely to increase the cover areas when using ruled surfaces to fit the surfaces. Meanwhile, we can trace another principal curvature direction to form the base curve. Accordingly, our cover generation step is driven by a smooth cross field approximating the principal curvature direction field. Central to the second stage is an optimization-based fitting to adaptively reduce the approximation error while keeping the ruled surfaces collision-free with the input shape. Our method's feasibility and practicability are demonstrated through various examples, including three physical manufacturing models using hot wire cutting.

Keywords: ruled surfaces, hot-wire cutting



Figure 1: Left: Hot Wire. Right: Diamond Wire. Images from <https://specialpatterns.com.au/> and <https://www.mactechonsite.com/>, respectively.

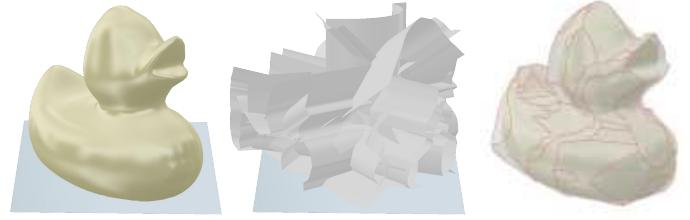


Figure 2: For an input shape with a fixed planar region (representing the workbench) (left), we generate a set of ruled surfaces (middle) which are collision-free with the input shape and fixed planar region so that cutting along these ruled surfaces removes as much material as possible (right).

1. Introduction

The subtractive fabrication process usually consists of two phases. One is rough machining, where the primary goal is to remove large amounts of material from a workpiece as quickly as possible [1, 2]. Another is using more precise techniques to shape the material for achieving the desired contours and dimensions necessary for the final product [3, 4]. Although rough machining is typically less precise, it is crucial for preparing the material for subsequent finer operations.

There are various ways of rough machining. One common method is to use regular wire to cut material, such as hot wire or diamond wire (Fig. 1). However, due to the limited motion space of robotic arms in diamond cutting and other methods, there is almost no optimization space. Therefore, we design our cutting path generation method for cutting methods with higher degrees of freedom, such as hot wire cutting. In this case, the surfaces formed by the movement of the regular wire are ruled surfaces. Given an input shape with a fixed planar region (representing the workbench), these ruled surfaces should satisfy the following requirements. First, each ruled surface is outside the input shape without any collision to ensure that the input shape is not destroyed after machining. Second, cutting along the constructed ruled surfaces leads to as much material removal as possible. Third, the number of ruled surfaces is small enough to improve the machining efficiency (Fig. 2).

In practice, users manually design such ruled surfaces; thus, it is a trial-and-error and time-consuming process. Our goal is to perform the design task automatically. However, this problem is rather challenging as computing ruled surfaces involves determining their numbers and placements under the overlap-free constraint. Generally, more material can be removed by introducing more ruled surfaces. Hence, the challenge is to find a favorable configuration (i.e., numbers and placements) of ruled surfaces.

Although [5] achieves globally collision-free ruled surfaces, their fitting error is high, and their method does not apply to complex models. Additionally, [6] explores hot wire cutting using curved lines; however, their method fails to guarantee collision-free and may result in overcutting. A potential approach for avoiding collisions is to scale up the input shape as the new input for [6]. However, finding an appropriate scale ratio requires manual tuning, which is tedious and may not generate results with low approximation error. Thus, it is unsuitable for practical machining processes. In general, the industry urgently needs an automatic algorithm that generates a small number of ruled surfaces to sufficiently approach the given shape without collisions.

In this paper, we propose a novel method to construct the desired ruled surfaces. We assume that the ruled surfaces are large enough to prevent the cutting tool from colliding with the material during the fabrication process. Our method is a two-step

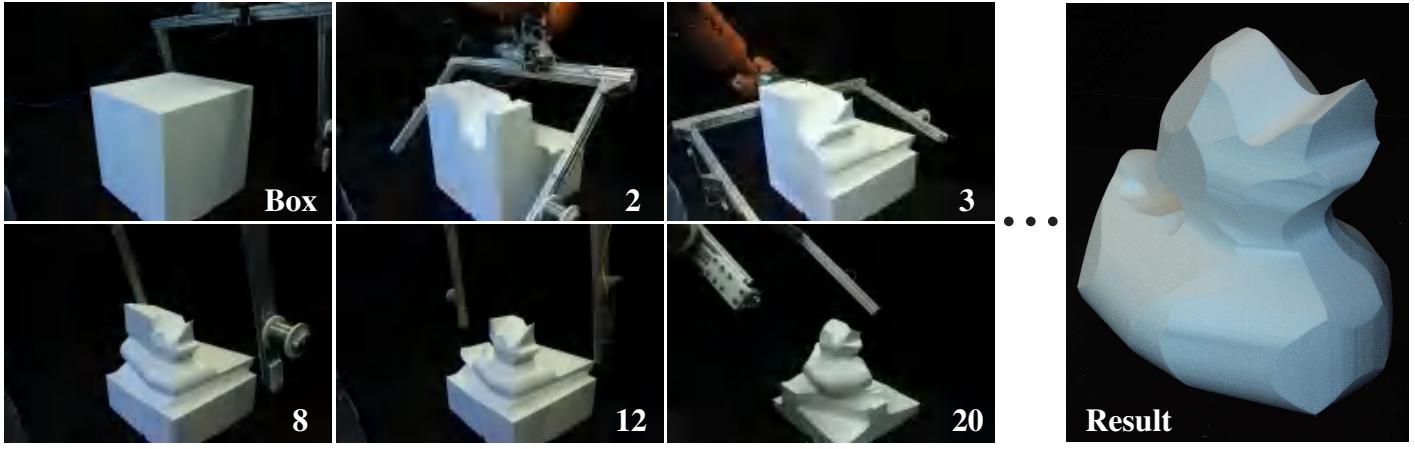


Figure 3: Given a shape fixed on the workbench, our algorithm automatically generates a small set of ruled surfaces that are collision-free with the input shape for rough machining. We achieve a trade-off between the number of ruled surfaces and the approximation error. Cutting along the optimized 23 ruled surfaces results in the rough machining result for the duck model. The number in the frame indicates the number of ruled surfaces that have been cut.

51 strategy that generates a small set of covers for the input and
 52 then fits cover patches using ruled surfaces without colliding with
 53 the input shape. However, developing a strategy to satisfy the
 54 requirements is challenging, and the reasons are twofold. First, the
 55 surface cover should facilitate the subsequent ruled surface fitting
 56 process, but this task has no well-known criterion. Second, ruled
 57 surface fitting with the collision-free constraint can be solved via
 58 a non-convex and nonlinear constrained optimization problem.
 59 However, it is often difficult to solve because the requirement that
 60 the ruled surface is sufficiently close to the input shape contradicts
 61 the collision-free constraint.

62 Our surface cover process is based on the following facts to
 63 overcome the first challenge. At a surface point, the normal
 64 change along one of the principal curvature directions is the fastest.
 65 Since the ruled surface can not be deformed along the rulings,
 66 setting the rulings and the principal curvature direction where the
 67 normal changes fastest orthogonal may increase the area that can
 68 be approximated well, called the cover area, thereby reducing
 69 the number of ruled surfaces. In this case, the other principal
 70 curvature direction can be used to guide the base curve directions.
 71 Consequently, we compute a smooth cross field close to the prin-
 72 cipal curvature direction field for surface cover to generate a set
 73 of cover patches. For the second challenge, we use a non-uniform
 74 B-spline surface (NURBS surface) to represent the ruled surface
 75 for fitting and present an interior point method to guarantee that
 76 the hard constraint is always met. To further reduce the approxi-
 77 mation error, we adaptively add control points at high error areas
 78 and assign them different step sizes. Moreover, using different
 79 step sizes allows us to abandon the barrier function, which speeds
 80 up the fitting process and reduces the approximation error.

81 In summary, our technical contributions are as follows:

- 82 • proposing an automatic algorithm to generate collision-free
 83 ruled surfaces for rough machining.
- 84 • offering an efficient and effective fitting algorithm to make
 85 the resulting ruled surfaces approach the input shape without
 86 collisions.

87 To the best of our knowledge, our system is the first to generate
 88 a set of ruled surfaces for rough machining automatically. It can
 89 be directly applied to digital fabrication, providing significant
 90 practical value to this field as the rough machining problem has
 91 been a long-standing and unresolved issue in the industry. We
 92 apply the developed tool to various shapes to demonstrate the

93 feasibility and effectiveness of our technique. Moreover, we
 94 fabricate three models physically using hot wire cutting along
 95 the generated ruled surfaces to show the system's usability and
 96 practicability.

2. Related Work

97 *Rough machining.* In subtractive manufacturing, the majority
 98 of rough machining work is based on a flat-bottomed cutter [7],
 99 which carves the external volume between the initial stock and
 100 the desired part shapes, following a predetermined tool path. As a
 101 result, rough machining typically accounts for 50% to 90% of the
 102 total machining time when a flat-bottomed cutter is used [8, 9].
 103 A variety of strategies have been developed to tackle the chal-
 104 lenge of planning roughing tool paths. These include the tradi-
 105 tional slicing-and-filling technique [10], roughing paths derived
 106 from tessellated models [11], specialized approaches for impeller
 107 blades [12, 3], a method of volume decomposition for 3-axis CNC
 108 rough machining of freeform 3D shapes [13], and a vector field
 109 based volume peeling for multi-axis machining [14]. Unlike the
 110 previous methods, we focus on rough machining using ruled sur-
 111 faces, employing a straight line cutter as the machining tool. In
 112 this area of research, [6] represents the most closely related work,
 113 where a deformed hot wire is applied to cut materials. However,
 114 their method may cut into the input model, violating the collision-
 115 free constraint, so it is unsuitable for rough machining. To our
 116 knowledge, this is the inaugural study to utilize ruled surfaces
 117 to conform to the desired part shape during the rough machining
 118 phase.

119 *Ruled surface-driven machining.* Beyond theoretical studies,
 120 many works have applied ruled surface fitting to practical mach-
 121 ining processes. [15] use the trajectory of a cylindrical tool
 122 to approximate ruled surfaces, conceptualizing ruled surfaces as
 123 curves on Plücker's quadric and determining the tool axis position
 124 using the striction line of the ruled surface. [16] also use
 125 cylindrical tools to fit ruled surfaces, employing an optimization
 126 approach to calculate the position of the tool head by the error
 127 between it and three guiding curves on the ruled surface, although
 128 their method lacks an automated initialization of guiding curves.
 129 Besides cylindrical tools, [17] approximate ruled surfaces using
 130 the motion of a conical tool, proposing an optimization method
 131 by sequentially optimizing the tool axis position via guiding rail

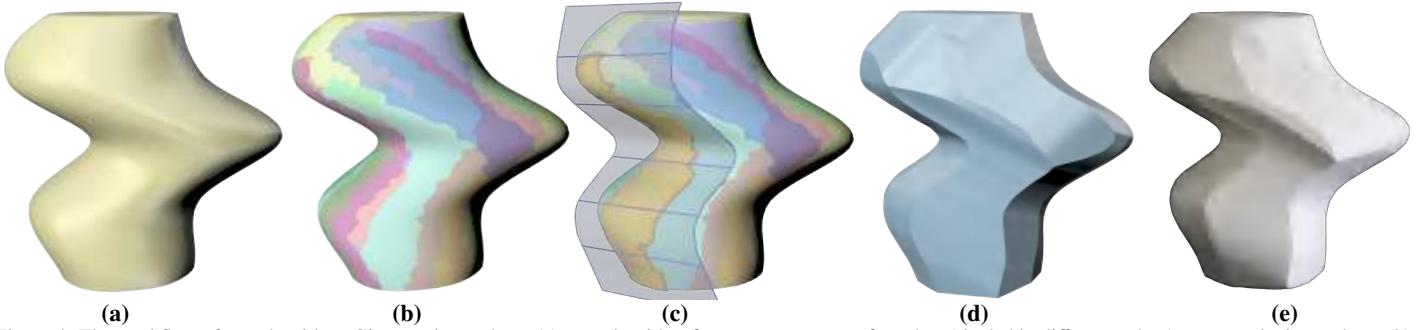


Figure 4: The workflow of our algorithm. Given an input shape (a), our algorithm first generates a set of patches (shaded in different colors) to cover the input shape (b). Then, we apply an adaptive fitting process to fit these patches with ruled surfaces (rendered as translucent surfaces) (c). We show the simulation result after cutting along the generated ruled surfaces (d) and the physically fabricated result with hot wire cutting (e).

movements. Recently, significant research has emerged on hot wire cutting. [18] develop a feed rate adjustment rule. It automatically controls the tool motion at feedrate-sensitive corners based on a bisection method, thus limiting the maximum machining errors and improving the machining accuracy of the tool path in the five-axis flank milling of ruled surfaces. [19] use hot wire cutting to manufacture minimal surfaces, with the cutting path moving along the principal curvature lines of the surface, cutting material from both sides. [5] propose a conservative algorithm for finding collision-free tangential cutting directions in hot wire cutting, based on which simple model cutting paths can be proposed.

Ruled surface fitting. Linear tools are often used in the fabrication process to cut materials for subsequent processing. The surface generated by machining using such tools is the ruled surface. So, there has been considerable work in the fields of architecture and manufacturing on fitting freeform surfaces with ruled surfaces. [20] adopt a subdivision-based method to approximate general freeform surfaces with ruled surfaces. They continuously subdivide the input surface along parametric directions until the user-defined error threshold is met. However, their method requires extensive subdivision steps if high accuracy is demanded. [21] propose a rationalization method for complex surfaces commonly found in architecture, using ruled surface strips to approximate freeform surfaces. Their approach involves optimizing the strips to ensure curvature continuity at the boundaries, providing a theoretical framework for ruled surface fitting. Based on the fact that any negatively curved freeform shape can be approximated by a smooth union of ruled surface strips, [22] achieve surface approximation by following a given surface's asymptotic curves with conical ruled surfaces' rulings. Additionally, [23] presents an approach using dynamic programming to compute piecewise ruled surface approximations of freeform surfaces. Different from them, we focus on automatically generating the global collision-free piecewise ruled surfaces for approximating a given shape.

3. Method

3.1. Problem and formulation

Inputs and goals. The input is a shape represented as a triangular mesh \mathcal{M} with a fixed planar region \mathbf{P}_{fix} (representing the work-bench). Without loss of generality, we assume that the input \mathcal{M} is inside a box \mathcal{B} with an edge length of 1, and the center of \mathcal{M} is at the origin. We aim to automatically generate a set of ruled surfaces $\mathcal{R} = \{\mathbf{R}_i\}$, which satisfy the following requirements:

- **Collision-free requirement:** each ruled surface \mathbf{R}_i does not collide with the input shape \mathcal{M} and is outside \mathcal{M} .

- **Material removal requirement:** each ruled surface \mathbf{R}_i is positioned as close to \mathcal{M} as possible to remove more material volume of \mathcal{B} .
- **Efficiency requirement:** the number of the ruled surfaces is small to improve machining efficiency.

The first requirement is a hard constraint.

Methodology. It is challenging to achieve a favorable configuration (i.e., numbers and placements) of ruled surfaces as the requirements are complex and mutually restrictive. In practice, we observe that if the collision-free requirement is not considered and \mathbf{R}_i is not required to be a ruled surface, constructing \mathcal{R} satisfying other requirements is just a shape cover problem. Then, after covering \mathcal{M} using a set of patches $\mathcal{P} = \{\mathbf{P}_i\}$, we can fit each patch \mathbf{P}_i using ruled surfaces under the collision-free constraint.

Pipeline. According to the cover-to-fit idea, we develop a two-stage approach: (1) cover \mathcal{M} with patches $\mathcal{P} = \{\mathbf{P}_i\}$ considering the property of the ruled surface, thereby facilitating the subsequent fitting stage (Section 4); and (2) robustly perform ruled surface fitting and adaptively add ruled surfaces to obtain a favorable tradeoff between the material removal requirement and the efficiency requirement while not intersecting \mathcal{M} (Section 5). Fig. 4 shows the workflow of our method.

4. Surface cover

Ruled surfaces. A ruled surface is a surface swept out by a straight line moving along a curve. The straight lines themselves are called rulings. The ruled surface has a parameterization of the form:

$$\mathbf{s}(u, v) = \mathbf{c}(u) + v\mathbf{t}(u), \quad (1)$$

where $\mathbf{c}(u)$ is the base curve and $\mathbf{t}(u)$ is a unit vector representing the direction of the ruling.

Tracing principal directions. We aim to have each patch in the surface cover set as close as possible to a ruled surface to facilitate subsequent fitting. Therefore, we need a criterion to assess the proximity of a surface to a ruled surface. Our key idea is to construct an approximation of the closest ruled surface to the target patch, as the exact closest one is hard to get. Since a ruled surface cannot bend along its ruling lines, setting rulings and the direction of maximum normal variation of the target patch (i.e., principal direction) orthogonal may be beneficial for making ruled surfaces fit larger areas. So, we trace the principal curvature directions and construct simulated rulings orthogonal to them. These traced principal curvature direction lines are then used as the simulated base curve to generate the ruled surfaces.

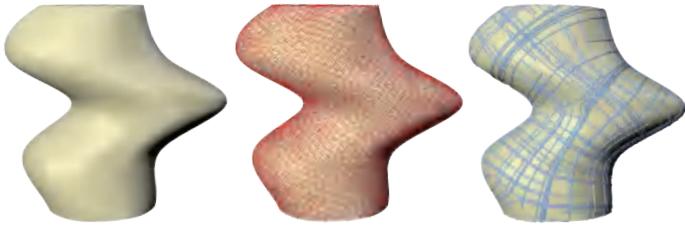


Figure 5: Given an input shape (a), we generate a smooth cross field (red crosses in (b)) that is used to find the direction lines L (blue lines in (c)).

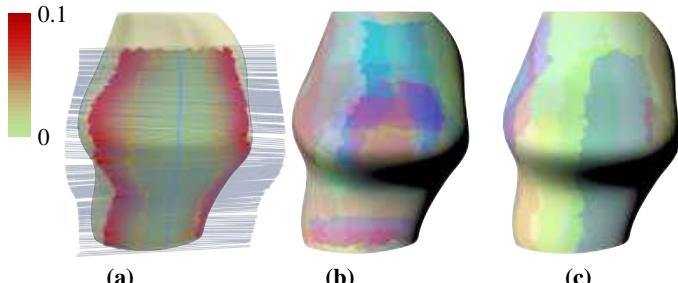


Figure 6: (a) We construct many mimic rulings (the translucent lines) on a direction line (the blue line) to simulate a ruled surface. The triangles whose distance to these mimic rulings below d_{cover} are collected to form a patch (enclosed by the orange line). The color (encoded by the color bar) on the shape indicates the distance from the shape to the rulings. (b) The redundant patch set $\mathcal{P}_{\text{redundant}}$. (c) The final patch set \mathcal{P} .

Smooth cross fields. In practice, it is difficult to determine the principal curvature directions in relatively flat places, and the principal curvature direction often contains many singularities. Thus, we compute a smooth cross field \mathcal{F} to approximate the principal curvature direction field. We follow [24, 25] to generate such a cross field by fixing the principal directions. Specifically, the cross field is encoded as a complex polynomial, and the conjugacy constraints are constructed based on the principal curvature directions. A smooth cross field satisfying the geometric constraints is iteratively generated by harmonically interpolating the polynomial coefficients and minimizing a smoothness energy.

Coarse-to-fine strategy. After constructing a smooth cross field, we propose a coarse-to-fine strategy to generate the cover patches \mathcal{P} . The coarse step mimics the properties of the ruled surface to generate many patches (Section 4.1). Then, a patch selection is conducted to determine the final patches (Section 4.2).

4.1. The coarse step

Mimicking base curves. We trace the cross field \mathcal{F} to form direction lines $\mathcal{L} = \{\mathbf{L}_i\}$ that mimic the base curve of ruled surfaces. We first randomly sample N_{seed} triangles as the starting points for tracing. To make the coarse step generate as many patches as possible for the subsequent refinement step, the two orthogonal directions given by the cross field \mathcal{F} on each triangle are all used for tracing to obtain enough direction lines. Given one direction \mathbf{d} , we propagate forward along the directions \mathbf{d} and $-\mathbf{d}$ starting from the center of \mathbf{f} with the same speed (see Fig. 5). The propagation process terminates if any of the following conditions are satisfied: (1) the propagation front reaches the one-ring neighborhood of a singularity of \mathcal{F} ; (2) the number of visited triangles exceeds $10\%N_{\text{tri}}$, where N_{tri} denotes the number of triangles of the input mesh. After propagation, we collect the ordered intersection points between the propagation rays and the mesh edges to form a direction line (denoted as $\mathbf{L} = \{\mathbf{p}_1, \dots, \mathbf{p}_n\}$), which is then put into \mathcal{L} .

We limit the length of the direction lines as a termination condition because a too-long direction line may produce a patch with a complex shape, making subsequent ruled surface fitting difficult and causing a high approximation error. The too-short direction lines do not impact our final results as we abandon them in the fine step. If the input mesh is irregular, we can use the average edge length or the diagonal length of the bounding box for the termination condition.

Mimicking rulings. Along each direction line $\mathbf{L} = \{\mathbf{p}_1, \dots, \mathbf{p}_n\}$, we construct many lines that mimic the rulings of the ruled surface. We select lines orthogonal to \mathbf{L} and the triangle normals of the input shape to mimic the rulings. The triangle where two adjacent points \mathbf{p}_i and \mathbf{p}_{i+1} locate is denoted as \mathbf{f}_i . We define a unit direction $\mathbf{r}_i = \mathbf{l}_i \times \mathbf{n}_i$, where $\mathbf{l}_i = (\mathbf{p}_{i+1} - \mathbf{p}_i)/\|\mathbf{p}_{i+1} - \mathbf{p}_i\|_2$ and \mathbf{n}_i is the normal direction of \mathbf{f}_i . On \mathbf{f}_i , we generate $\lfloor \|\mathbf{p}_{i+1} - \mathbf{p}_i\|_2/\delta \rfloor + 1$ lines of length β in the direction \mathbf{r}_i and centered at \mathbf{p}_i^k , where δ indicates a small distance and $\mathbf{p}_i^k = \mathbf{p}_i + k\delta\mathbf{l}_i$, $k = 0, \dots, \lfloor \|\mathbf{p}_{i+1} - \mathbf{p}_i\|_2/\delta \rfloor + 1$, respectively. In our experiments, we set $\delta = 0.25d_{\text{ae}}$ and $\beta = d_{\text{bb}}$, where d_{bb} denotes the length of the bounding box's diagonal and d_{ae} is the average edge length of the input shape.

Patch construction. To simulate the following ruled surface fitting process under the global collision-free constraint, the generated line at point \mathbf{p}_i^k is translated along the face normal by $d_{\text{translation}}$. The ruling lines without translations may intersect with the input model, thus violating collision-free constraints. As a result, the cover area may be reduced, increasing the number of patches. Therefore, we translate the ruling lines by a specific distance to balance the number of patches and the approximation error. It enlarges the cover area of each patch, thereby reducing the number of patches.

For each direction line \mathbf{L} , we first sample N_{sample} points on each generated line. Then, we perform the closest point projection for each sample point to the mesh, where the closest triangle is visited. If the minimum projection distance for a visited triangle is greater than d_{cover} , the triangle is marked as unvisited. Finally, we perform the flood fill algorithm from the triangles, where \mathbf{L} passes through, to collect the visited triangles for forming a patch (Fig. 6 (a)). Since the mimic rulings represent the cutting tool, overcutting occurs once they collide with the input shape. So, we abandon the patch in which the mimic rulings intersect with the input shape as it violates the collision-free requirement. We run these steps for all direction lines to generate a redundant patch set $\mathcal{P}_{\text{redundant}}$ (Fig. 6 (b)). In our experiments, we set $N_{\text{sample}} = 200$, $d_{\text{translation}} = 1.5\%d_{\text{bb}}$, and $d_{\text{cover}} = 2.5\%d_{\text{bb}}$.

4.2. The fine step

Set cover. This step aims to select as few patches as possible from $\mathcal{P}_{\text{redundant}}$ to cover the input shape due to the efficiency requirement (see Fig. 6 (c)). We first perform a set cover algorithm [26] on $\mathcal{P}_{\text{redundant}}$ to achieve a minimum set cover \mathcal{P}_{set} .

Greedy refinement. The result of the set cover algorithm may include patches that highly overlap with others. To further reduce the number of patches, we develop a greedy algorithm to eliminate redundant patches. First, we sort the patches in \mathcal{P}_{set} in descending order based on their area, denoted as $\{\mathbf{P}_{\text{sort},i}\}$. The resulting patch set $\mathcal{P} = \{\mathbf{P}_{\text{sort},1}\}$ is initialized by pushing the patch with the largest area into the set. Then, we iterate the sorted patches sequentially from a large to a small area. If a patch has 90% of its area covered by the patches in \mathcal{P} , we skip it; otherwise, we add it into \mathcal{P} . This process terminates until all the sorted patches are traversed or the union of the areas of all patches in \mathcal{P} exceeds 99% of the area of

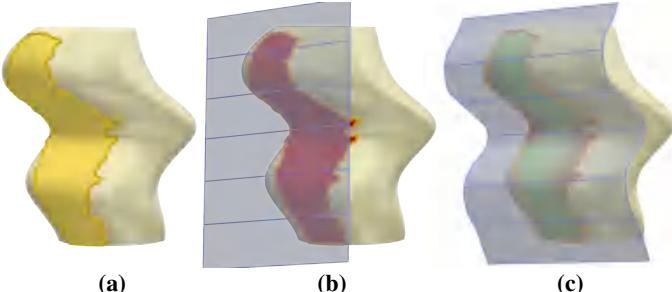


Figure 7: The fitting process. (a) The input shape \mathcal{M} (the yellow model) and one patch $\mathbf{P} \in \mathcal{P}$ (the orange part) to be fitted. (b) The initial ruled surface. (c) The fitting result. The color on the patch in (b, c) indicates the distance from the patch point to the ruled surface with the color bar in Fig 6.

313 the input shape. For each triangle that is not covered, we assign it
 314 to its one adjacent patch in \mathcal{P} if the patch has the smallest sum of
 315 the angle difference between its normal and the patch's average
 316 normal and the distance between its center and the patch center.

5. Ruled surface generation

5.1. Fitting algorithm

319 **Problem overview.** The inputs of the fitting process include the
 320 input shape \mathcal{M} and one patch $\mathbf{P} \in \mathcal{P}$, which has a direction line \mathbf{L}
 321 and mimicked rulings. We aim to generate a ruled surface \mathbf{R} to
 322 approach \mathbf{P} while not colliding with \mathcal{M} (see Fig. 7).

323 **Representation.** We use a NUBS surface to represent the ruled
 324 surface \mathbf{R} . A B-spline curve in three-dimensional (3D) space is
 325 specified by a knot vector $\mathbf{u} = \{u_0, u_1, \dots, u_m\}$ with $u_i \leq u_{i+1}$,
 326 basis functions of degree $p \in \mathbb{N}$ ($p = 3$ in our experiments), and
 327 control points $\{\mathbf{p}_i \in \mathbb{R}^3, i = 0, 1, \dots, n\}$, with the relationship
 328 $m = p + n + 1$:

$$\mathbf{c}(u) = \sum_{i=0}^n B_{i,p}(u) \mathbf{p}_i, \quad \forall u \in [u_0, u_m], \quad (2)$$

329 where the basis function $B_{i,p}(u)$ is:

$$B_{i,0}(u) = \begin{cases} 1 & \text{if } u_i \leq u \leq u_{i+1} \\ 0 & \text{otherwise} \end{cases}, \quad (3)$$

$$B_{i,p}(u) = \frac{u - u_i}{u_{i+p} - u_i} B_{i,p-1}(u) + \frac{u_{i+p+1} - u}{u_{i+p+1} - u_{i+1}} B_{i+1,p-1}(u) \quad \text{if } p \geq 1. \quad (4)$$

331 Given two B-spline curves $\mathbf{c}_0, \mathbf{c}_1$ with identical knot vectors,
 332 the linear interpolation between corresponding points $\mathbf{c}_0(u), \mathbf{c}_1(u)$
 333 sharing the same parameter u is performed to define the ruled
 334 surface \mathbf{R} , which has a parameterization of the form:

$$\begin{aligned} \mathbf{R}(u, v) &= (1 - v)\mathbf{c}_0(u) + v\mathbf{c}_1(u) \\ &= \sum_{i=0}^n B_{i,p}(u) \{(1 - v)\mathbf{p}_{i,0} + v\mathbf{p}_{i,1}\}, \end{aligned} \quad (5)$$

335 where $\mathbf{p}_{i,j}$ ($j = 0$ or 1) is the control point of \mathbf{c}_j and $v \in [0, 1]$. We
 336 collect the control points of the two B-spline curves into a set \mathcal{C}
 337 and the knot vectors into a set \mathcal{U} .

Algorithm overview. The shape of the ruled surface is determined
 338 by these two sets \mathcal{C} and \mathcal{U} . Accordingly, we apply an iterative
 339 strategy, which alternates in each iteration a pass to optimize the
 340 positions of control points and a pass to add control points (i.e.,
 341 subdivide the knot vector) to generate the resulting ruled surface.
 342 First, our algorithm optimizes the position of control points to
 343 approach the input shape under the collision-free constraint (Sec-
 344 tion 5.1.1). The number of control points and the knot vector are
 345 fixed during this step. Then, we add new control points at the
 346 regions of the ruled surface with high approximation errors (Sec-
 347 tion 5.1.2). These two steps are alternately performed. Our ini-
 348 tialization and termination conditions are shown in Section 5.1.3.
 349 Since adding new control points corresponds to subdividing the
 350 knot vector, it can be considered that we only optimize \mathcal{C} . We
 351 denote the ruled surface determined by \mathcal{C} as $\mathbf{R}_\mathcal{C}$.

5.1.1. Updating positions of control points

Optimization problem. The fitting process is formulated as the
 353 following optimization problem:

$$\begin{aligned} \arg \min_{\mathcal{C}} \quad & E(\mathcal{C}) \\ \text{s.t.} \quad & \mathbf{R}_\mathcal{C} \cap \mathcal{M} = \emptyset, \end{aligned} \quad (6)$$

356 where the optimization variable is \mathcal{C} and $E(\mathcal{C})$ denotes the ob-
 357 jective function.

Objective function. Our objective function contains two parts:

$$E(\mathcal{C}) = E_{\text{error}}(\mathcal{C}) + \omega E_{\text{smooth}}(\mathcal{C}). \quad (7)$$

359 Here, ω ($\omega = 1$ in our experiments) is a weight to balance the
 360 approximation error and the smoothness energy. $E_{\text{error}}(\mathcal{C})$ is:

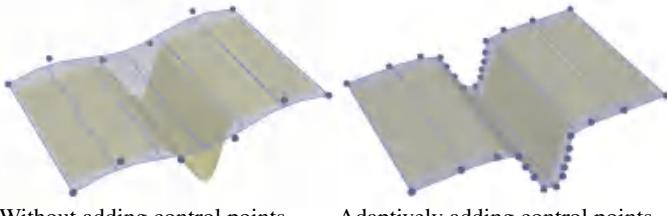
$$E_{\text{error}}(\mathcal{C}) = \sum_{\mathbf{v}_k \in \mathbf{P}} \|\mathbf{v}_k - \mathbf{R}_\mathcal{C}(u_k, v_k)\|_2^2, \quad (8)$$

361 where \mathbf{v}_k denotes a vertex within \mathbf{P} and $\mathbf{R}_\mathcal{C}(u_k, v_k)$ represents the
 362 closest projection of \mathbf{v}_k onto the ruled surface $\mathbf{R}_\mathcal{C}(u, v)$. $E_{\text{smooth}}(\mathcal{C})$
 363 is:

$$\begin{aligned} E_{\text{smooth}}(\mathcal{C}) &= E_{\text{smooth},u}(\mathcal{C}) + \omega_{uv} E_{\text{smooth},v}(\mathcal{C}), \\ E_{\text{smooth},u}(\mathcal{C}) &= \sum_{i=0}^{n-2} \sum_{j=0}^1 \|\mathbf{p}_{i,j} - 2\mathbf{p}_{i+1,j} + \mathbf{p}_{i+2,j}\|_2^2, \\ E_{\text{smooth},v}(\mathcal{C}) &= \sum_{i=0}^{n-1} \|(\mathbf{p}_{i,0} - \mathbf{p}_{i,1}) - (\mathbf{p}_{i+1,0} - \mathbf{p}_{i+1,1})\|_2^2, \end{aligned} \quad (9)$$

364 where the parameter ω_{uv} balances the smoothness terms along
 365 u-direction and v-direction. $E_{\text{smooth},u}$ and $E_{\text{smooth},v}$ control the
 366 smoothness of control points in the u-direction and v-direction,
 367 respectively. We will discuss other choices of $E_{\text{smooth},u}$ in Section
 368 5.1.3. ω_{uv} is set to 5 in the experiments.

Collision-free constraint. According to the convex hull property
 369 of the NUBS surface, we perform subdivisions to ensure that the
 370 ruled surface avoids collision with the input mesh \mathcal{M} . Specifically,
 371 subdividing the ruled surfaces in the guide curve direction k times
 372 and in the ruling direction l times yields sub-surfaces $\{\mathbf{R}_{i,j}\}$. If
 373 the convex hull of the control points of each $\mathbf{R}_{i,j}$ avoids collision
 374 with any mesh face $\mathbf{f} \in \mathcal{M}$, the ruled surface \mathbf{R} does not intersect
 375 with the input shape \mathcal{M} . It is essential to select appropriate values
 376 for k and l to ensure that each subdivided subsurface $\mathbf{R}_{i,j}$ closely
 377 corresponds in scale to the faces of the input mesh. To enhance
 378 the precision of collision detection while optimizing the use of
 379 computational resources, we set $k=8$ and $l=7$ in our experiments.



Without adding control points Adaptively adding control points
Figure 8: Updating knot vectors to add control points.

Solver overview. We stick to the following principle to guarantee that the constraint is always met: initialize a constraint-satisfied ruled surface and then optimize it without violating the constraint. To optimize the problem (6) concerning two types of variables, i.e., C and the auxiliary variables $\mathbf{R}_C(u_k, v_k)$ in (8), we use the local-global solver. In the local step, we compute $\mathbf{R}_C(u_k, v_k)$ while fixing the control points C , i.e., fixing the ruled surface. The global step fixes (u_k, v_k) for \mathbf{v}_k to update C .

Local step: updating the closest point. Given a vertex \mathbf{v}_k within \mathbf{P} , we compute the closest point $\mathbf{R}_C(u_k, v_k)$ using an open-source toolkit called Open Cascade Technology [27]. Then, we establish a correspondence between the vertex \mathbf{v}_k and the parameters (u_k, v_k) . In the global step of updating the control points C , we assume that this correspondence is fixed.

Global step: updating positions of the control points. Given the correspondence between each \mathbf{v}_k and (u_k, v_k) , the objective function is quadratic concerning the positions of the control points C . We can easily minimize the quadratic function to obtain a solution C^* if the collision-free constraint is not considered. However, the ruled surface with C^* may intersect with \mathcal{M} .

A trivial solution for problem (6) is to use barrier functions to reformulate it as an unconstrained optimization problem. However, in our experiments, we observe that the barrier function not only prevents the ruled surface \mathbf{R}_C from fitting the patch \mathbf{P} well but also slows down the optimization process (see Fig. 15 (c) and more discussions in Section 6). So, we choose to abandon the barrier function and propose an adaptive line search to ensure our results meet the collision-free constraint. Specifically, we find that once there is a part of \mathbf{R}_C close enough to \mathbf{P} , it may cause zero step size in the line search step, which prevents the approximation error from further reducing. To avoid this, we assign the same step sizes for a pair of control points $(\mathbf{p}_{i,0}, \mathbf{p}_{i,1})$, whereas adopting different step sizes for different pairs due to the locality of B-spline basis functions.

The descent direction for each control point is defined as the difference between the solved and the original positions. We update one pair of control points $(\mathbf{p}_{i,0}, \mathbf{p}_{i,1})$ while fixing other control points every time as follows:

1. We perform the Continuous Collision Detection (CCD) algorithm between the convex hull (generated by subdivision) and the input mesh to find the maximum allowable step size t_{\max} , ensuring that the collisions are avoided.
2. We perform a standard Armijo backtracking algorithm to determine the step size, which is initialized as $0.5t_{\max}$.

Each pair of control points corresponds to a segment of the ruled surface according to their influence range. The greater the E_{error} of the corresponding segment of a pair of control points, the earlier the pair of control points is updated.

5.1.2. Adding control points

The number of control points C has a significant impact on the fitting ability of the ruled surface \mathbf{R}_C . To further reduce the approximation error in the high-error areas, we increase the number of control points at the corresponding segments of the ruled surface \mathbf{R}_C (Fig. 8). The new control points are added as follows. First, we sort the segments of the knot vector according to their approximation error. Then, we find the segment with as large error as possible while satisfying the following condition to avoid generating too short segments through subdivision: the length of the ruled surface strip corresponding to this segment is greater than a threshold γ (we set $\gamma = 5d_{\text{ae}}$). More details about how to add control points are shown in the supplementary material. This process increases the number of control points, thereby reducing each control point's influence region to better accommodate complex shapes.

5.1.3. Other details

Initialization. Since the interior-point method is employed, each step produces a collision-free ruled surface, making a collision-free initialization essential. We choose a plane as the initial ruled surface with the following steps: (1) a B-spline curve with 6 control points is fitted to the given principal curvature lines; (2) a ruled surface is reconstructed from the fitted B-spline curve by optimizing the ruling to be orthogonal to the B-spline curve and the input shape's face normal; (3) the principal component analysis (PCA) method is used to reduce the dimensionality of the reconstructed ruled surface to a plane; (4) we translate the plane along its normal direction to a collision-free position as the initial plane. The details are shown in the supplementary material.

Termination conditions. The algorithm stops when the average distance E_{average} , i.e., the average value of Euclidean distances from vertices in \mathbf{P} to the ruled surface, is less than a given threshold ϵ or when the number of iterations exceeds the specified maximum iteration number N_{iter} . In our experiments, we set $N_{\text{iter}} = 10$ and $\epsilon = 2\%d_{\text{bb}}$.

Testing multiple initializations. If $E_{\text{average}} > \epsilon$ after fitting for some patch, we rotate the initialization plane and run the fitting process to further reduce E_{average} . The rotation operation is uniformly performed N_{rotate} times both along the plane's normal and rulings. Then, we choose the result that has the smallest E_{average} as our fitting result. In our experiments, we set $N_{\text{rotate}} = 15$.

In practice, our initializations can produce fitting results that satisfy $E_{\text{average}} \leq \epsilon$ for most cover patches, eliminating the need for rotating sampling.

Smoothness term along u-direction. [28] provided a smoothness energy of a B-spline curve as follows:

$$E_{\text{smooth}}^{\text{curve}}(\mathbf{c}(u)) = \alpha \int_{u_0}^{u_m} \|c'(u)\|_2^2 du + (1-\alpha) \int_{u_0}^{u_m} \|c''(u)\|_2^2 du. \quad (10)$$

where $0 \leq \alpha \leq 1$ and $c(u)$ is the B-spline curve. This smoothness energy can also be used in (9) as the u-direction smoothness term. However, according to (4), the denominator of $c'(u)$ and $c''(u)$ contains the knot vector's interval length. Thus, subdividing the knot vector may cause the denominator to approach zero so that the value of $c'(u)$ and $c''(u)$ explodes. To avoid this issue, we finally choose to use $E_{\text{smooth}, u}$ in (9), a trivial smoothness term, to measure the smoothness of the ruled surfaces.

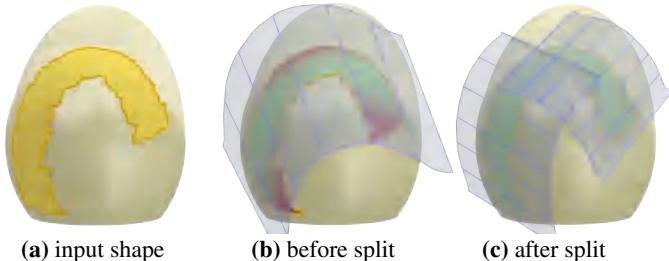


Figure 9: We perform an adaptive split method to the patches have high approximation error. The color on the cover patch in (b, c) encodes the distance from the point of the cover patch to the ruled surface using the color bar in Fig 6.

483 5.2. Adaptive splitting and fitting

484 **Motivation.** After performing the above fitting process, we ob-
 485 serve that E_{average} of some patches is still larger than ϵ (see Fig. 9
 486 (b)). We think the reason is that the shape of these patches is too
 487 complex requiring the direction of the rulings to change drastically,
 488 so it is hard to use one ruled surface to fit it well. Thus, we
 489 introduce an adaptive splitting process to reduce the approxima-
 490 tion error. Its core is to segment the patches with large errors into
 491 smaller patches for further fitting.

492 **Segmentation.** We use the elegant spectral mesh seg-
 493 mentation algorithm [29] to split a patch \mathbf{P} , given its fitting ruled surface \mathbf{R} .
 494 The segmentation algorithm is based on the following energy:

$$E_{\text{seg}} = \sum_{\mathbf{e} \in \mathcal{E}_I} E_{\text{rotation}}(\mathbf{e}) + \omega_{\text{seg}} E_{\text{distance}}(\mathbf{e}), \quad (11)$$

495 where \mathcal{E}_I denotes the interior edge set of \mathbf{P} . E_{rotation} indicates the
 496 rotation difference between normals, E_{distance} denotes the distance
 497 error, and ω_{seg} is a positive parameter. E_{seg} is defined on the
 498 interior edges of \mathbf{P} , which is beneficial for segmentation.

499 We define E_{rotation} as follows:

$$E_{\text{rotation}}(\mathbf{e}) = \|\mathbf{q}_{\text{patch}}(\mathbf{e}) - \mathbf{q}_{\text{ruled}}(\mathbf{e})\|_2^2. \quad (12)$$

500 Here, we denote the two adjacent faces of \mathbf{e} as \mathbf{f}_1 and \mathbf{f}_2 . The
 501 quaternion $\mathbf{q}_{\text{patch}}(\mathbf{e})$ represents the rotation between the two
 502 normal directions of triangles \mathbf{f}_1 and \mathbf{f}_2 . We perform the closest point
 503 projection for the triangle centers of \mathbf{f}_1 and \mathbf{f}_2 to the ruled surface
 504 \mathbf{R} to get two points. The rotation of the normal directions at these
 505 two points is denoted by $\mathbf{q}_{\text{ruled}}(\mathbf{e})$. E_{rotation} indicates the rotation
 506 difference between normals and estimates the second derivative
 507 of the approximation error, preventing the error from fluctuating
 508 greatly, which is crucial for fine machining. $E_{\text{distance}}(\mathbf{e})$ is the sum
 509 of the closest distance from the four vertices of triangles \mathbf{f}_1 and \mathbf{f}_2
 510 to the ruled surface \mathbf{R} . It is four times E_{average} of the two neighbor
 511 triangles of \mathbf{e} , measuring the local approximation error.

512 **The splitting process.** We run the splitting process as follows:

- 513 1. Select all the patches with E_{average} exceeding ϵ to form \mathcal{A} .
- 514 2. Pop one patch \mathbf{P} from \mathcal{A} and perform the segmentation on it
 515 to generate two sub-patches \mathbf{P}_1 and \mathbf{P}_2 .
- 516 3. Use ruled surfaces to fit \mathbf{P}_1 and \mathbf{P}_2 .
- 517 4. Add the sub-patch whose E_{average} exceeding ϵ to \mathcal{A} .
- 518 5. If \mathcal{A} is empty, stop the process; otherwise, go to Step (2).

519 In our experiments, we perform a maximum of three split opera-
 520 tions on each initial patch to avoid excessive patches. Actually,
 521 each initial patch is usually split only once. In practice, the
 522 connectivity of triangles can influence the energy defined on edges.
 523 In our experiments, the meshes are regular. For irregular meshes,
 524 subdividing the inputs is an effective strategy.

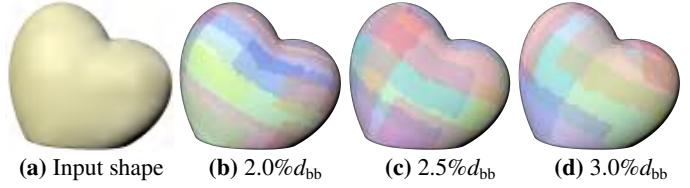


Figure 10: The cover parameter d_{cover} . From left to right, the number of cover patches progressively decreases, i.e., (21, 15, 10).

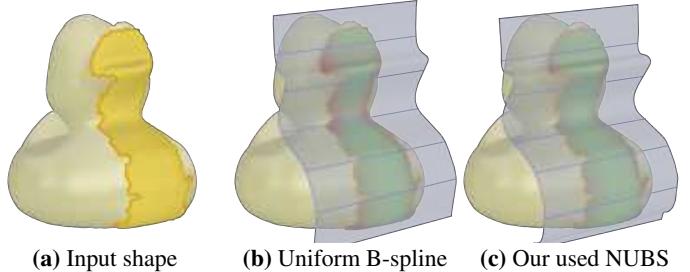


Figure 11: Different representations of the ruled surfaces. The number of control points of the uniform B-spline is the same as our resulting NUBS surface. The color on the cover patch in (b, c) encodes the distance from the point of the cover patch to the ruled surface using the color bar in Fig 6.

525 6. Experiments and evaluations

526 We have tested our algorithm on various models to evaluate its
 527 performance, including three physical fabrication examples. Our
 528 algorithm was implemented in C++, and the experiments were
 529 executed on a desktop PC with a 3.80 GHz Intel Core i7-10700
 530 and 32 GB of memory. Tables 1 and 2 summarize the statistics
 531 of our results. Unless otherwise specified, we use the color bar in
 532 Fig 6 to show the approximation error or distance.

533 6.1. Algorithm evaluations

534 **The cover parameter d_{cover} .** In the surface cover process, we use
 535 the parameter d_{cover} to control the area of each patch. We test
 536 different d_{cover} s to show its effect, as shown in Fig. 10. If d_{cover}
 537 is small (e.g., $d_{\text{cover}} = 2.0\%d_{\text{bb}}$ in Fig. 10 (b)), the area of each
 538 patch is small and the number of patches is great. Instead, a large
 539 d_{cover} (e.g., $d_{\text{cover}} = 3.0\%d_{\text{bb}}$ in Fig. 10 (d)) leads to a large area
 540 of each patch, making it challenging to fit each of them with low
 541 approximation errors. To achieve a favorable trade-off between
 542 the number of patches and the difficulty of the fitting process, we
 543 set $d_{\text{cover}} = 2.5\%d_{\text{bb}}$ in our experiments. We list the parameters of
 544 our algorithm in the appendix.

545 **Representations of ruled surfaces.** Uniform B-splines and NUBS
 546 can also be used to represent ruled surfaces. Compared to uniform
 547 B-splines, non-uniform B-splines (NUBS) are more effective in
 548 handling regions with high curvature variation. The key advan-
 549 tage lies in their non-uniform knot vectors. In uniform B-splines,
 550 knots are evenly spaced, which results in a fixed influence range
 551 for each control point. This fails to adaptively increase control
 552 point density in regions with rapid curvature variation, which lead
 553 to significant approximation errors (see Fig 11 (b)). In contrast,
 554 the knot vectors of NUBS can be adjusted according to curva-
 555 ture. This allows for more accurate approximation of complex
 556 geometries (see Fig 11 (c)). Hence, we use NUBS to represent
 557 ruled surfaces and apply an adaptive fitting method to handle this
 558 problem.

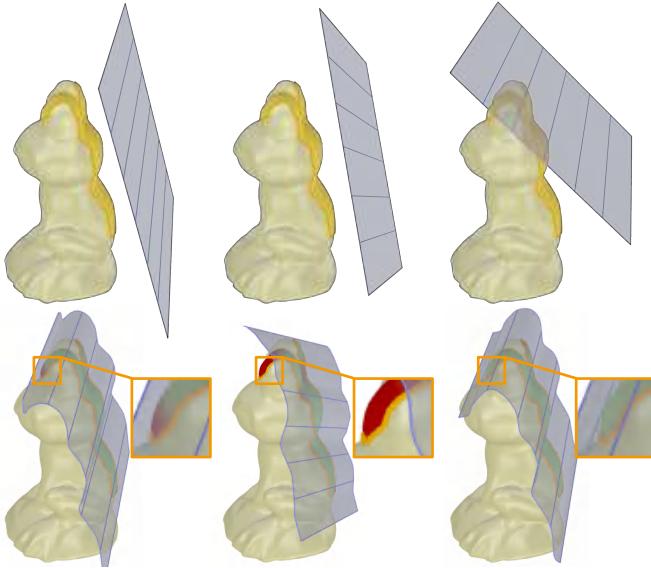


Figure 12: Different initializations for the fitting process. We show the initializations in the upper and the fitting results in the bottom.

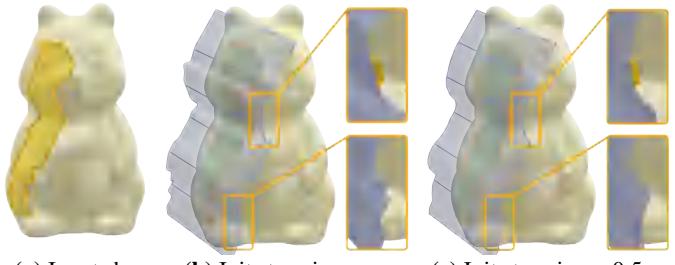


Figure 13: Different initial step sizes during the line search.

559 *Initializations for the fitting process.* The initial plane in the
 560 complex fitting process has a significant impact on the results. We test
 561 three initializations for the ruled surface fitting in Fig. 12. Thus,
 562 if the fitting error is large, we apply the rotational sampling to
 563 generate a suitable initialization (see the details in Section 5.1.3).

564 *Parameters in fitting.* There are two parameters in the fitting
 565 process that affect the shape of the final ruled surface. The first one
 566 is the initial step size in the line search for updating control points
 567 (Fig. 13). In our experiments, we set it $0.5t_{\max}$ where t_{\max} is the
 568 maximum allowable step size satisfying the collision constraint.
 569 This is because decreasing the movement of control points in
 570 each iteration helps improve the fairing of the ruled surfaces.
 571 The second one is the minimum length of the ruled surface that
 572 can be subdivided in the updating knot vectors step. As shown in
 573 Fig. 14, if we do not limit the length of the ruled surface, it may be
 574 excessively subdivided, leading to ruled surfaces with fluctuation.
 575 (Fig. 14 (b)). Therefore, we set the minimum length to be d_{ae} in
 576 our experiments.

577 *Other fitting solutions.* [30] provide a trivial solution for solving
 578 the fitting problem. They employ a barrier function to ensure
 579 the ruled surface is always outside the input shape and utilize
 580 Newton's method to update the control points:

$$clog(x) = \begin{cases} -\frac{(x-x_0)^2}{x} \log\left(\frac{x}{x_0}\right) & \text{if } 0 < x \leq x_0 \\ 0 & \text{elsewhere} \end{cases} \quad (13)$$

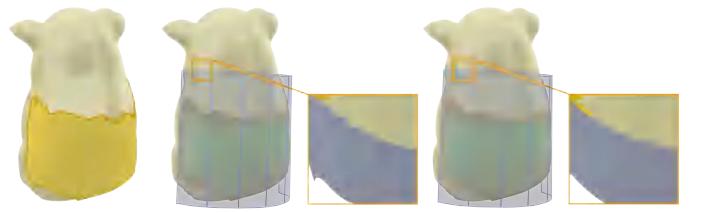


Figure 14: Different minimum lengths for the ruled surfaces to be subdivided.

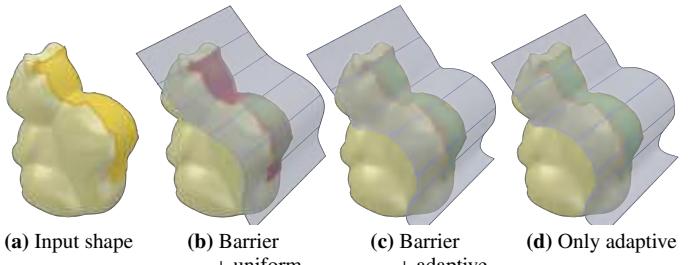


Figure 15: Different fitting strategies. (a) The input shape. (b) Using barrier functions and uniform step sizes ($E_{\text{average}} = 0.00412$, time = 9 minutes). (c) Using barrier function and different step sizes for control points ($E_{\text{average}} = 0.00046$, time = 30 minutes). (d) Only using different step sizes for control points ($E_{\text{average}} = 0.00022$, time = 6 minutes). The color on the cover patch in (b, c, d) encodes the distance from the point of the cover patch to the ruled surface using the color bar in Fig 6.

The variable x represents the distance from a sample point on the input model to the ruled surface and x_0 is a user-defined parameter. They then use this barrier function to reformulate the problem with hard constraints as an unconstrained one. It is easy to solve but still has problems. Using a uniform step size in the line search may lead to certain parts of the target surface not fitting well (see Fig. 15 (b)). This is because, to meet the no-collision constraint, areas that fit well contribute zero step sizes in the line search process. Since all control points share the same step size, the part of the ruled surface corresponding to the not well-fitted areas will stagnate, which results in great approximation error between the ruled surface and the target surface. To address this issue, we assign different step sizes to control points to achieve adaptive fitting with the barrier function (Fig. 15 (c)).

However, there are still two problems. First, the barrier function produces opposite descent directions in areas close to the target surface, trying to prevent the ruled surface from colliding with it. Second, computing different step sizes for each control point involves more computational complexity, which increases the method's running time.

In our experiments, we observe that if we abandon the barrier function and only use line search to ensure the two surfaces do not collide, we can quickly achieve a sufficiently low energy result (Fig. 15 (d)). Although the final results may not necessarily be local optimal solutions, nor can we achieve first-order optimality conditions since the descent direction we calculate is not the correct one, they are good enough for the fabrication process. Since the energy continuously decreases, we can set the step size as the termination condition. Almost all results satisfy the criterion of energy being below a threshold because the shape of the target patches is not complex. Since integrating strict collision-free constraints into the energy formulation is challenging and time-consuming, we use adaptive step sizes without the barrier function for efficient fitting with low approximation errors for practicality.



Figure 16: We succeed in generating a few ruled surfaces for rough machining for nine models.

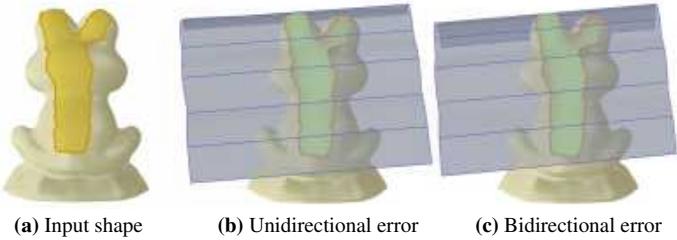


Figure 17: *Unidirectional error vs bidirectional error.* (a) The input shape. (b) Unidirectional error ($E_{\text{average}} = 0.00053$, time = 9 minutes). (c) Bidirectional error ($E_{\text{average}} = 0.00045$, time = 15 minutes).

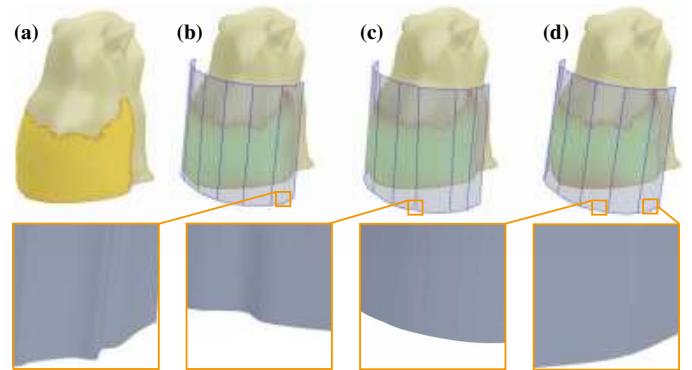


Figure 18: *E_{smooth} term.* (a) The input shape. (b) Only $E_{\text{smooth},u}$ ($E_{\text{average}} = 0.00071$). (c) Only $E_{\text{smooth},v}$ ($E_{\text{average}} = 0.00062$). (d) Ours ($E_{\text{average}} = 0.00056$).

616 *One-directional error vs bidirectional error.* The E_{error} in Equ. 7
617 can also be measured by a bidirectional error. Fig. 17 shows
618 the comparison between the bidirectional error result and ours.
619 The experiment shows that the fitting error based on the one-
620 directional distance is comparable to that of the bidirectional distance,
621 but the computational efficiency is reduced when using the
622 bidirectional error. To achieve a good trade-off between approxi-
623 mation error and computational efficiency, we use One-directional
624 error as E_{error} in Equ. 7.

625 *Smoothness term.* The smoothness of a B-spline surface (e.g., its
626 continuity C^k) is determined by its knot vectors and polynomial
627 degree. However, this inherent smoothness only ensures local conti-
628 nuity in the parametric domain and does not directly guarantee geo-
629 metric smoothness. When fitting complex shapes—particularly in
630 regions with high curvature variation—relying solely on the
631 parametric continuity of B-splines may lead to geometric artifacts
632 such as folds or sharp transitions. To address this, additional
633 smoothness terms (e.g., $E_{\text{smooth},u}$ and $E_{\text{smooth},v}$) are introduced to
634 enforce global surface smoothness from a geometric perspective,
635 rather than relying purely on parametric continuity (see Fig. 18
636 (b)). The smoothness term $E_{\text{smooth},v}$ is used to control the smooth-
637 ness of the ruled surface along the v-direction. It is computed

638 based on the differences between corresponding control points
639 along the v-direction. This helps suppress undesirable oscillations
640 and promotes smoothness of the ruled surface (see Fig. 18 (c)).

641 *The fitting error vs time.* In our experiments, the number of con-
642 trol points can affect the fitting accuracy and efficiency. Fig. 19
643 shows the three results with different numbers of control points.
644 For fairness, we control the number of final total control points by
645 adjusting parameters ϵ and γ . As shown in Fig. 19, using fewer
646 control points reduces computation time but results in higher fit-
647 ting error. Conversely, using more control points yields accuracy
648 comparable to our method, but results in a significantly higher
649 computational cost. Based on this trade-off, we adopt $\epsilon = 3\%d_{\text{bb}}$
650 and $\gamma = 5d_{\text{ae}}$ as our default parameter settings.

651 *Comparison with practitioner’s results.* We invite a practitioner
652 from the field of industrial robot processing to generate the cut-
653 ting paths with Rhinoceros 3D software [31]. We compare the

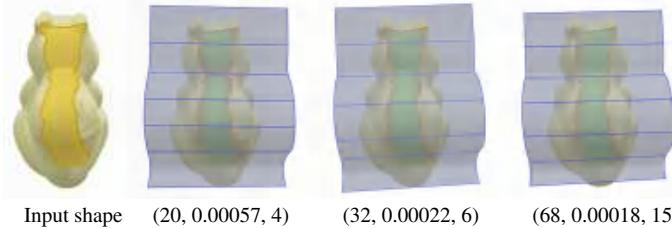


Figure 19: The fitting error vs time. The parameter used in the second column are $\epsilon = 5\%d_{bb}$ and $\gamma = 6d_{ae}$. The third column shows the result produced using our method with its default parameter setting, which are $\epsilon = 3\%d_{bb}$ and $\gamma = 5d_{ae}$. The parameter used in the last column are $\epsilon = 1\%d_{bb}$ and $\gamma = 4d_{ae}$. The text below each result indicates the number of control points, $E_{average}$, and time consumption (mins), respectively.

practitioner's results with three examples (see Fig. 20). To better evaluate our algorithm, the practitioner is asked to produce results under two termination conditions: (1) the $E_{average}$ less than that of ours and (2) the number of cuts that are the same as ours. The experiment shows that the approximation error and the number of cuts of the practitioner's results are slightly better than that of ours in all three examples. However, the practitioner consumes much longer time than our method. Even in the example with the smallest gap (the first and last in the third row in Fig 20), the practitioner's time is 2.6 times ours (2.53 and 0.96 hours respectively). Our automated algorithm does not need manual labor while producing results with comparable quality to manual results. Thus, our method has significant practical value in digital fabrication.

More examples. We show nine models in Fig. 16. For these shapes, we successfully generate small sets of ruled surfaces for rough machining. Fig. 21 shows a simulation fabrication process, where cutting along the generated ruled surfaces leads to the rough machining result. The other three examples physically fabricated by the hot wire cutting are shown in Fig. 23. These simulation and fabrication results demonstrate the robustness and effectiveness of our algorithm.

Timings. The surface cover and the ruled surface fitting took about 1.32 minutes and 1.17 hours for the duck model with 6434 vertices and 23 ruled surfaces in Fig. 3. The fitting process occupies most of our method as we need to find proper step sizes for different control points so that the ruled surfaces satisfy the collision-free constraint and are close enough to the target patch. Among this process, the CCD for finding the maximum allowable step size is the most time-consuming. In addition, the collision-free initialization time in the ruled surface fitting is almost negligible. Although our algorithm is not fast enough, it is more time-saving compared to manual labor and is a fully automated algorithm. It produces results with low approximation error and a small number of cuts, thus having a significant practical value. We summarize the time consumption for all models in the paper in Table 1.

6.2. Physical fabrication

Hardware setup. Our physical setup includes a KUKA® Quantec KR120 R2700 Robot extra HA, which has a 6-degree-of-freedom arm with a working range of 2696mm (Fig. 22). In front of it, we have placed a workbench and a polystyrene foam cube with an edge length of 600mm, designated as the object to be cut. The hot wire cutting tool consists of a frame, flange, right-angle connectors, insulation blocks, springs, wire clamps, a coil of hot wire, and an adjustable DC power source. The frame is constructed by connecting one 6060 alloy profile and two 3060 aluminum

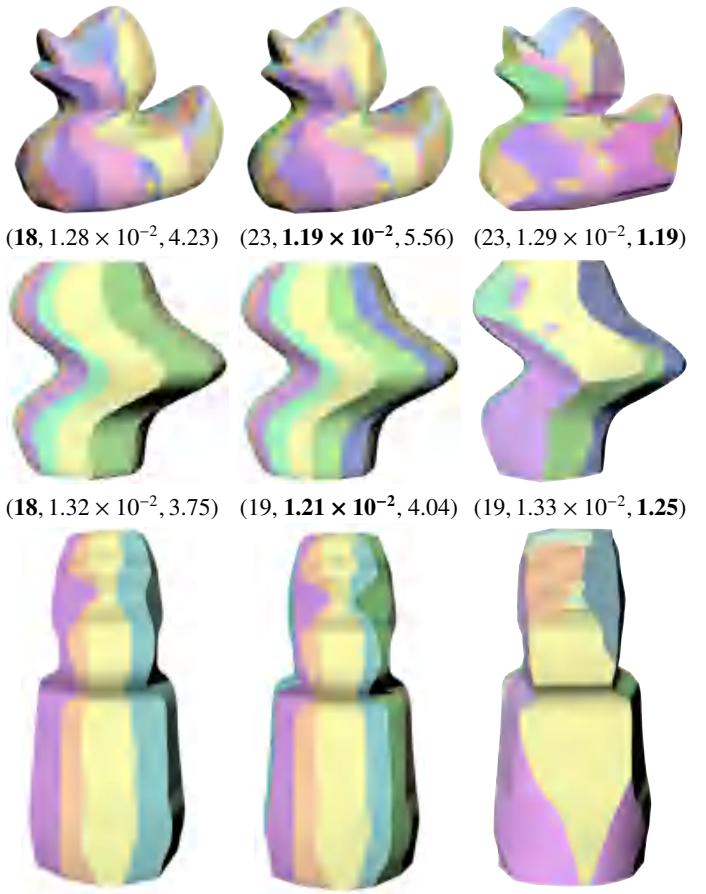


Figure 20: Comparison with practitioner's results. The first column shows the practitioner's results which are produced when their $E_{average}$ is less than that of ours. The second column includes the practitioner's results with the same number of cuts as ours. Our results are in the last column. The text below each result indicates the number of cuts, average error, and time consumption (hours), respectively.

Models	Number of cuts	Our result tol 2% d_{bb}	Our result tol 3% d_{bb}	d_{result}^{avg}	d_{bb}	t_{cover} (h)	$t_{fitting}$ (h)
Fig. 3 Duck	23	94.88%	98.85%	1.29×10^{-2}	1.63	2.19×10^{-2}	1.17
Fig. 4 Artwork	19	94.34%	98.94%	1.33×10^{-2}	1.58	2.11×10^{-2}	1.13
Fig. 16 Drill	17	92.89%	98.68%	0.95×10^{-2}	1.06	2.08×10^{-2}	1.60
Fig. 16 Foot	22	96.58%	99.88%	1.01×10^{-2}	1.13	2.53×10^{-2}	1.92
Fig. 16 Frog	38	92.29%	97.84%	1.01×10^{-2}	1.43	4.64×10^{-2}	3.03
Fig. 16 Head	31	96.80%	99.33%	0.98×10^{-2}	1.51	3.47×10^{-2}	2.65
Fig. 16 Heart	25	99.76%	100.00%	0.93×10^{-2}	1.67	2.39×10^{-2}	0.92
Fig. 16 Pig	27	93.13%	98.57%	1.31×10^{-2}	1.47	2.92×10^{-2}	2.35
Fig. 16 Squirrel	42	93.28%	98.01%	1.29×10^{-2}	1.48	5.36×10^{-2}	3.40
Fig. 16 Venus	31	96.62%	99.66%	1.02×10^{-2}	1.21	2.61×10^{-2}	1.68
Fig. 16 Barrel	19	99.95%	100.00%	1.14×10^{-2}	1.73	1.56×10^{-2}	0.47
Fig. 21 Moai	14	87.75%	95.09%	1.29×10^{-2}	1.21	1.69×10^{-2}	0.98

Table 1: Overview of our simulation shapes. We measure the distance from our simulation shape to the input shape to define the error. d_{result}^{avg} is average distance. The third (or fourth) column shows the ratio of the area within the specified error threshold to the total area of the simulation shape. t_{cover} and $t_{fitting}$ denote the time consumed by the surface cover process and the fitting process, respectively. These times are measured in hours.

profiles using right-angle connectors. The flange is mounted in the middle of the 6060 aluminum profile; two insulation blocks are mounted at the ends of the two 3060 aluminum profiles—one insulation block is equipped with a spring, and the other with a wire clamp. The hot wire coil is fixed on one of the 3060 aluminum profiles, and the hot wire is drawn from the coil. One end of the hot wire passes through the wire clamp and is fixed on the spring of the other 3060 aluminum profile. The wire clamp and



Figure 21: A simulation fabrication process of one example. The left three columns are the first, second, and fourteenth cuts. The last two columns show the input shape and the simulation shape generated by the simulated hot wire cutting.



Figure 22: Hardware Setup. Robotic arm product model: KUKA® Quantec KR120 R2700 extra HA Robot.

Models	Number of cuts	Our result tol $\pm 2\% d_{bb}$	Our result tol $\pm 3\% d_{bb}$	d_{fab}^{avg}	d_{bb}
Fig. 3 Duck	23	98.96%	99.80%	1.00×10^{-2}	1.63
Fig. 4 Artwork	19	100.00%	100.00%	0.49×10^{-2}	1.58
Fig. 21 Moai	14	92.26%	98.69%	0.85×10^{-2}	1.21

Table 2: Overview of the fabrication results. We use the one-sided distance from the fabrication shape to our simulation shape to measure the fabrication error. d_{fab}^{avg} is mean error in the fabrication. The third (or fourth) column shows the ratio of the area within the specified error threshold to the total area of the fabrication surface.

708 spring are equipped with lead-out wires connected to the positive
709 and negative terminals of the adjustable DC power source. In
710 our experiments, we set the power source current to 2.0 amperes
711 and the voltage to 20.0 volts, heating the hot wire to 250 degrees
712 Celsius. The movement speed of a point on the hot wire is set
713 to 2mm/s, ensuring that the foam melts before contacting the hot
714 wire, thereby avoiding the impact of resistance.

715 **Motion trajectory.** We use NUBS to represent the motion trajectory
716 of the robotic arm. The motivation is that smooth B-spline
717 surfaces typically result in lower machining errors than polylines
718 in practical manufacturing. This is because the robotic arm must
719 pause or significantly decelerate at the turning points of polyline

720 paths or accelerate abruptly when changing direction. Such actions
721 can cause excessive dwelling at sharp corners, resulting in
722 unwanted material melting or inducing deformation due to rapid
723 directional changes, thus negatively affecting fabrication accuracy.

724 **Machining order.** We assume that the ruled surfaces representing
725 the cutting surfaces are sufficiently wide. Hence, the machining
726 order has almost no impact on the fabrication results. Besides,
727 mature algorithms (e.g., KUKA—prc) have been developed in the
728 industry to connect multiple cutting paths. Therefore, we study
729 how to generate each cutting path rather than the complete path.

730 Besides, to ensure the model does not fall off the material in
731 advance during the cutting process, we manually select the last
732 cutting path in the software.

733 **Entry and exit cut.** In the physical fabrication process, we need to
734 specify additional cutting paths for entry and exit cuts, which are
735 not discussed in previous work [5, 6]. This is because the cutting
736 tool (a hot wire in our experiments) has a high temperature, and if
737 it is close to the surface of the input shape, the material may be
738 burned and damaged. In our algorithm, since the ruled surface we
739 used to fit the patch is large enough, the redundancy part of the
740 ruled surface naturally provides a path for entry and exit cuts. We
741 only need to set the hot wire cut along the ruled surfaces, and the
742 material will be protected from being burned.

743 **Fabrication constraints.** Since the material is placed on a platform
744 in front of the cutting machine, we need to consider whether
745 the robotic arm collides with the platform and the material during
746 the fabrication process. To ensure that our results meet this fabri-
747 cation constraint, we add a bottom plane for the input shape. The
748 method is as follows. For each input shape, we manually add a
749 rectangular plane as the bottom plane at an appropriate position.
750 Then, this bottom plane is added to the collision detection in our
751 algorithm to generate a set of ruled surfaces that do not intersect
752 with the input shape and the bottom plane. So, at this point, the
753 cutting path generated based on our results will naturally not cause
754 collisions between the robotic arm and the platform. The col-
755 lisions between the robotic arm and the material can be eliminated
756 by appropriately setting the size of the unprocessed material \mathcal{B} . When
757 \mathcal{B} is small enough, there exists a way to place the robotic
758 arm to avoid collisions during cutting. In our experiments, we set
759 the edge length of \mathcal{B} as half the length of the hot wire.

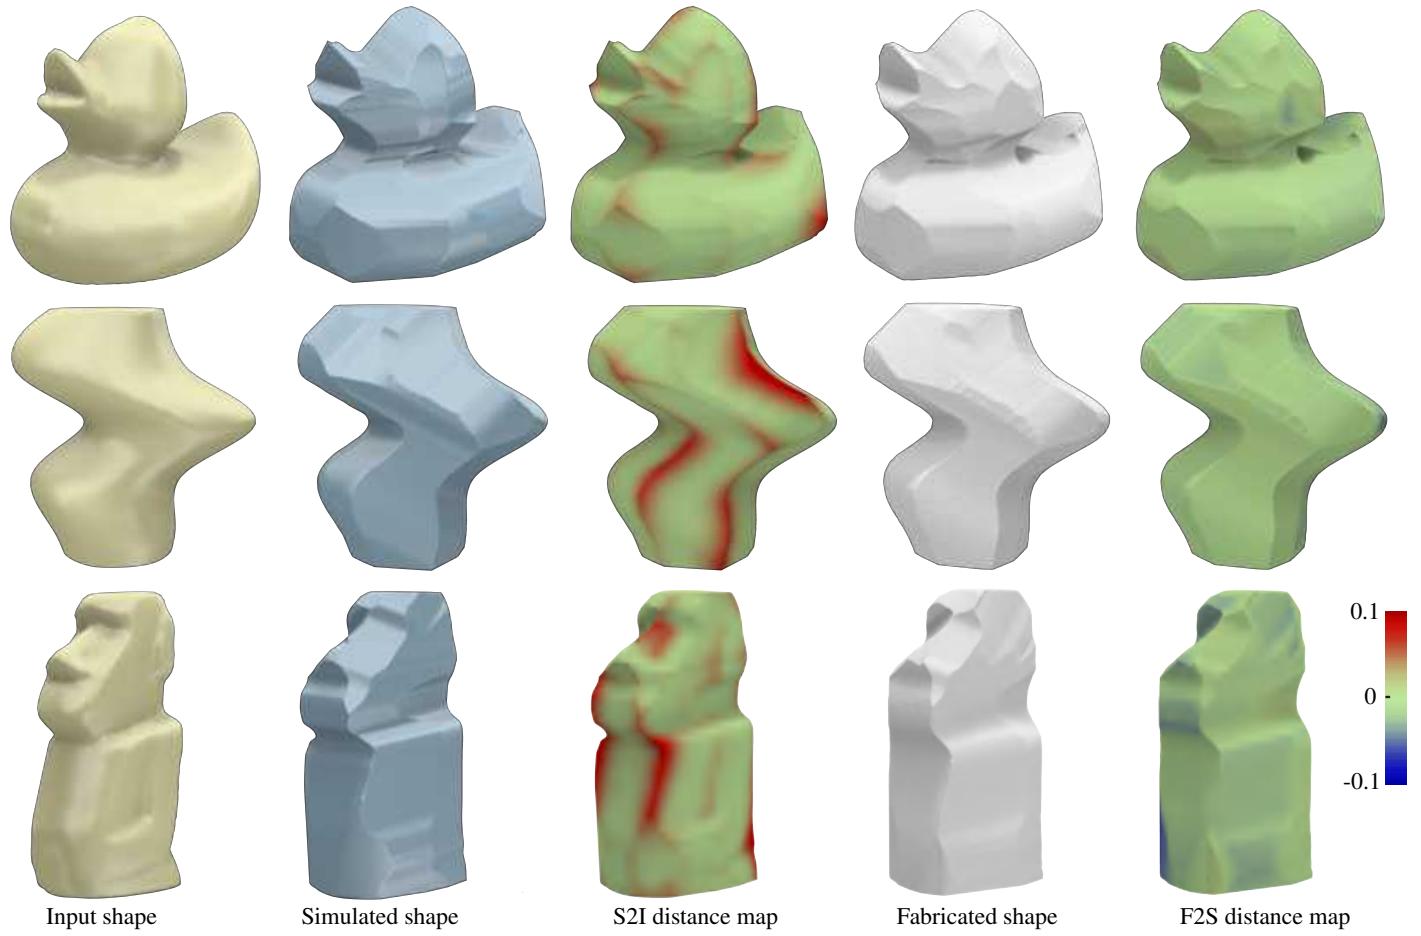


Figure 23: Simulation and fabrication errors. The term “S2I” refers to simulation shapes to input shapes, and “F2S” refers to fabrication shapes to simulation shapes. The distance map is color-coded by the color bar.

60 **Analysis.** We test our algorithm on 12 different shapes. Table. 1
 61 shows the analysis of these results. For these complex shapes,
 62 our algorithm can achieve results with small approximation errors
 63 using about 30 cuts (see Table. 1 second column). Only the
 64 Frog and Squirrel models need to be cut by 38 and 42 times.
 65 The distances from our simulation shape to the input shape in
 66 most areas are below 2% d_{bb} and almost the distances of all parts
 67 are below 3% d_{bb} (Table. 1 third and fourth columns). We also
 68 evaluate the accuracy of the fabrication results (see Table. 2). We
 69 use a 3D scanner to reconstruct these models and then compute the
 70 distance from the fabrication results to our simulation results. The
 71 two errors are small and comparable. This shows our algorithm is
 72 enough to use in the physical fabrication process.

7. Conclusion and Discussions

774 We propose a novel method to compute a small set of ruled
 775 surfaces for rough machining. Each ruled surface is strictly out-
 776 side the input shape. A cover-to-fitting strategy is developed to
 777 achieve a good trade-off between the number of ruled surfaces and
 778 the approximation error under the intersection-free constraints.
 779 Specifically, our adaptive fitting strategy makes ruled surfaces
 780 fully approximate the input surface without using a barrier func-
 781 tion. We also fabricated three models using hot wire cutting to
 782 demonstrate the feasibility and practicability of our algorithm.
 783 To the best of our knowledge, our work is the first to offer an
 784 automated solution to this problem. Although our algorithm has

785 a slow overall speed, we believe it has practical value due to its
 786 fully automated nature and sufficiently small approximation error
 787 of the results.

788 **Choice of initial ruled surfaces.** Constructing the initial ruled
 789 surfaces based on contour lines rather than using planar surfaces
 790 provides a more geometrically reasonable initial surface, which
 791 can guide the optimization away from a local optimum and ac-
 792 celerate convergence. However, contour lines can only generate
 793 limited types of ruled surfaces, such as cylindrical and planar
 794 surfaces, so they cannot handle conical surfaces well. How to
 795 solve this difficulty and apply contour lines to our algorithm is a
 796 worthwhile direction for future research.

797 **Approximation error.** For an arbitrary model, we can not ensure
 798 that the approximation error between the final result and the input
 799 shape is always less than a specified threshold. Since the purpose
 800 of rough machining is to efficiently remove as much material as
 801 possible, rather than obtaining a precise model, we only carry
 802 out a limited number of patch splits to reduce the approximation
 803 error.

804 **Joint patch optimization.** The patches are optimized inde-
 805 pendently rather than jointly, meaning that the partition bound-
 806aries cannot be adjusted during optimization. This restriction prevents
 807 the correction of large errors at patch intersections and limits the
 808 overall surface fitting quality. Although joint optimization can
 809 improve fitting results, it significantly increases the cost of fitting.

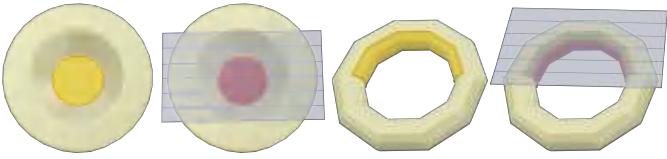


Figure 24: Left: an example of cutting concave shape. Right: an example of cutting high-genus shape.

810 **This is primarily due to the need to simultaneously optimize thousands of control points while performing collision detection. The time required to solve such a large-scale optimization problem becomes practically unacceptable. Moreover, joint optimization would necessitate dynamically adjusting the fitting region. How to properly assign fitting regions to each ruled surface while ensuring full coverage of the input shape remains a challenging problem.**

811 **Concave shapes.** Regular hot wire cutting machines cannot handle some kinds of concave shapes because the tool will collide with the input shape when cutting the concave part with a large depth. So, to maintain the collision-free constraint, the distance from this concave part to the ruled surface should be larger than the depth, which makes it unacceptable for rough machining (Fig. 24 left). This shortcoming is inherent in this kind of machining method. Using curves instead of straight lines as the hot wire can, to a certain extent, solve this problem [6]. It is worth exploring in this direction in the future.

812 **None-zero genus shapes.** Similar to concave shapes, there are also areas on none-zero genus shapes that can not be cut by the hot wire (Fig. 24 right). A feasible solution is to partition the input shape into several genus-zero parts and process them separately. However, it may compromise the integrity of our results. Another simple way is to manually make holes in the input shape and then reassemble the hot wire to cut through the holes. However, this method is undoubtedly very cumbersome. In the future, we aim to improve our algorithm to be able to deal with high-genus shapes.

813 **The type of the input shape.** Since our method utilizes the principal curvature direction field to guide the generation of the ruled surfaces, it is unsuitable for handling shapes with noisy surfaces. In this case, the principal curvature direction field is not smooth, which makes the lengths of the mimic base curves short. This decreases the area each path can cover and then increases the number of ruled surfaces in the final results. A feasible solution is to construct a cage for this input shape and run our method to process the cage. In addition, shapes with many planes, such as CAD models, are also not suitable for our method to handle. This is because the principal curvature directions of a plane can be arbitrary, while our method just traces one of them. Our method does not optimize the principal curvature direction to ensure no collision between the ruled surface constructed from the principal curvature direction and the input shape. Designing a dedicated algorithm is a better choice for these special cases, as our algorithm is more inclined to handle models with general shapes.

853 References

- 854 [1] Z. Yin, Rough and finish tool-path generation for nc machining of freeform
855 surfaces based on a multiresolution method, Computer-Aided Design 36 (12)
856 (2004) 1231–1239.
- 857 [2] G. Vickers, Z. Dong, H. Li, Optimal rough machining of sculptured parts on
858 a cnc milling machine, Journal of Engineering for Industry, Transactions of
859 the ASME 115 (4) (1993) 424–432.
- 860 [3] H.-T. Young, L.-C. Chuang, K. Gerschwiler, S. Kamps, A five-axis rough
861 machining approach for a centrifugal impeller, The International Journal of
862 Advanced Manufacturing Technology 23 (2004) 233–239.
- 863 [4] S. Tao, K.-L. Ting, Unified rough cutting tool path generation for sculptured
864 surface machining, International Journal of Production Research 39 (13)
865 (2001) 2973–2989.
- 866 [5] B. van Sosin, M. Bartoň, G. Elber, Accessibility for line-cutting in freeform
867 surfaces, Comput. Aided Des. 114 (2019) 202–214.
- 868 [6] S. Duenser, R. Poranne, B. Thomaszewski, S. Coros, Robocut: Hot-wire
869 cutting with robot-controlled flexible rods, ACM Trans. Graph. 39 (4) (2020)
870 98–1.
- 871 [7] Y. Hu, W. Tse, Y. Chen, Z. Zhou, Tool-path planning for rough machining
872 of a cavity by layer-shape analysis, The International Journal of Advanced
873 Manufacturing Technology 14 (5) (1998) 321–329.
- 874 [8] E.-Y. Heo, D.-W. Kim, B.-H. Kim, D.-K. Jang, F. F. Chen, Efficient rough-
875 cut plan for machining an impeller with a 5-axis nc machine, International
876 Journal of Computer Integrated Manufacturing 21 (8) (2008) 971–983.
- 877 [9] J. E. Petrzalka, Geometric process planning in rough machining (2009).
- 878 [10] M. Liang, S. Ahamed, B. Van Den Berg, A step based tool path generation
879 system for rough machining of planar surfaces, Computers in Industry 32 (2)
880 (1996) 219–231.
- 881 [11] M. Balasubramaniam, P. Laxmiprasad, S. Sarma, Z. Shaikh, Generating 5-
882 axis nc roughing paths directly from a tessellated representation, Computer-
883 Aided Design 32 (4) (2000) 261–277.
- 884 [12] Y. Chen, G. Mei, M. Lv, J. Gao, Tool path generation for efficient rough-cut
885 machining of ruled surface impellers, in: 2014 International Conference on
886 Mechatronics and Control (ICMC), 2014, pp. 599–602.
- 887 [13] A. Mahdavi-Amiri, F. Yu, H. Zhao, A. Schulz, H. Zhang, Vdac: volume
888 decompose-and-carve for subtractive manufacturing, ACM Transactions on
889 Graphics (TOG) 39 (6) (2020) 1–15.
- 890 [14] N. Dutta, T. Zhang, G. Fang, I. E. Yigit, C. C. Wang, Vector field based
891 volume peeling for multi-axis machining, in: International Design Engineering
892 Technical Conferences and Computers and Information in Engineering
893 Conference, Vol. 87295, American Society of Mechanical Engineers, 2023,
894 p. V002T02A025.
- 895 [15] K. Sprott, B. Ravani, Cylindrical milling of ruled surfaces, The International
896 Journal of Advanced Manufacturing Technology 38 (2008) 649–656.
- 897 [16] H. Gong, L.-X. Cao, J. Liu, Improved positioning of cylindrical cutter for
898 flank milling ruled surfaces, Comput. Aided Des. 37 (12) (2005) 1205–1213.
- 899 [17] C. Li, S. Bedi, S. Mann, The International Journal of Advanced Manufacturing
900 Technology 29 (2006) 1115–1124.
- 901 [18] C. Chu, W. Huang, Y. Hsu, Machining accuracy improvement in five-axis
902 flank milling of ruled surfaces, International Journal of Machine Tools and
903 Manufacture 48 (7–8) (2008) 914–921.
- 904 [19] H. Hua, T. Jia, Wire cut of double-sided minimal surfaces, The Visual
905 Computer 34 (6) (2018) 985–995.
- 906 [20] L. Piegl, W. Tiller, The nurbs book, Springer-Verlag, New York (1995).
- 907 [21] S. Flöry, H. Pottmann, Ruled surfaces for rationalization and design in
908 architecture, LIFE in: formation. On responsive information and variations
909 in architecture (2010) 103–109.
- 910 [22] S. Flöry, Y. Nagai, F. Isovranu, H. Pottmann, J. Wallner, Ruled free forms,
911 in: Advances in architectural geometry 2012, Springer, 2013, pp. 57–66.
- 912 [23] C. C. Wang, G. Elber, Multi-dimensional dynamic programming in ruled
913 surface fitting, Comput. Aided Des. 51 (2014) 39–49.
- 914 [24] O. Diamanti, A. Vaxman, D. Panozzo, O. Sorkine-Hornung, Designing n-
915 polyvector fields with complex polynomials, Comput. Graph. Forum 33 (5)
916 (2014) 1–11.
- 917 [25] D. Bommes, H. Zimmer, L. Kobbelt, Mixed-integer quadrangulation, ACM
918 Trans. Graph. 28 (3) (2009).
- 919 [26] W. Lin, F. Ma, Z. Su, Q. Zhang, C. Li, Z. Lü, Weighting-based parallel
920 local search for optimal camera placement and unicost set covering, in:
921 Proceedings of the 2020 Genetic and Evolutionary Computation Conference
922 Companion, 2020, pp. 3–4.
- 923 [27] O. C. S.A.S., Opencascade technology, <http://www.opencascade.com>
(2020).
- 924 [28] Y. Liu, H. Yang, W. Wang, Reconstructing b-spline curves from point clouds—
925 a tangential flow approach using least squares minimization, in: International
926 Conference on Shape Modeling and Applications 2005 (SMI’05), IEEE,
927 2005, pp. 4–12.
- 928 [29] W. Tong, X. Yang, M. Pan, F. Chen, Spectral mesh segmentation via ℓ_0
929 gradient minimization, IEEE T. Vis. Comput. Gr. 26 (4) (2018) 1807–1820.
- 930 [30] R. Ni, T. Schneider, D. Panozzo, Z. Pan, X. Gao, Robust & asymptotically
931 locally optimal uav-trajectory generation based on spline subdivision, in:
932 2021 IEEE International Conference on Robotics and Automation (ICRA),
933

Table C.3: Summary of parameters used in our method

Parameter	Symbol	Value
difference of two principal curvatures	χ	10
diagonal length of the bounding box	d_{bb}	-
average edge length	d_{ae}	-
length of a ruling line	β	d_{bb}
a small distance	δ	$0.25d_{ae}$
number of sampling triangles	N_{seed}	200
number of sampling points on each direction line \mathbf{L}	N_{sample}	200
distance of translation	$d_{translation}$	$1.5\%d_{bb}$
minimum projection distance	d_{cover}	$2.5\%d_{bb}$
the length of the ruled surface strip	γ	$5d_{ae}$
number of subdivisions in the guide curve direction	k	8
number of subdivisions in the ruling direction	l	7
maximum iteration number	N_{iter}	10
number of rotation	N_{rotate}	15
approximation error bound	ϵ	$2\%d_{bb}$
smoothness parameter	ω_{uv}	5
segmentation parameter	ω_{seg}	0.5

937 Appendix A. Adding control points details

938 We use the following criterion to add control points:

$$939 N_{new} = \begin{cases} 5 & \text{if } E_{interval} > 3\epsilon \text{ and } L_{interval} > 2\gamma \\ 0 & \text{if } E_{interval} < \epsilon \text{ or } L_{interval} < \gamma \\ 3 & \text{otherwise} \end{cases}, \quad (A.1)$$

940 where N_{new} is the number of new points, $E_{interval}$ is this interval's
 941 approximation error E_{error} and $L_{interval}$ is this interval's length. The
 942 positions of the new points are determined by dividing the interval
 943 into $N_{point} + 1$ equal segments, with the maximum value of N_{point}
 limited to 5.

944 Appendix B. Initialization details

945 **Step 1:** The B -spline $\mathbf{c}(u)$ fitting problem can be formulated as
 946 follows:

$$947 \arg \min_C E_{dist} = \sum_{k=0}^K \|\mathbf{v}_k - \mathbf{c}(u_k)\|_2^2 / |\mathcal{V}|, \quad (B.1)$$

948 where $\mathcal{V} = \{\mathbf{v}_k\}$ are the sampling points on the fitted principal
 949 curvature line \mathbf{L} , $\mathbf{c}(u_k)$ is the closest point to v_k and C is the control
 950 points of $\mathbf{c}(u)$. The initialization of C is set to some samples in \mathcal{V} .
 951 The first and last control points are constrained to be the start and
 952 end points of \mathbf{L} .

953 This problem is easy to solve because it is a least square prob-
 954 lem by solving the normal equations. We also adaptively add
 955 control points to make the fitting result better. The iteration stops
 956 when E_{dist} is less than α . In our experiments, we set α to the
 957 average edge length of the input mesh.

958 **Step 2:** We uniformly sample the knot vector of $\mathbf{c}(u)$ to obtain
 959 $\{u_m\}$. Then we construct $\{\mathbf{l}_m\}$ passing through $\{u_m\}$. The direction
 960 of $\{\mathbf{l}_m\}$ is orthogonal to $\mathbf{c}(u)$'s tangent vector and the input shape's
 961 face normal. We set the length of $\{\mathbf{l}_m\}$ to d_{bb} and the center of $\{\mathbf{l}_m\}$
 962 at $\{u_m\}$.

963 Then we generate a ruled surface \mathbf{s} to fit $\{\mathbf{l}_m\}$. We use the
 964 same representation of ruled surfaces in Formula (5). So we
 965 denote $\mathbf{c}_0(u)$ and $\mathbf{c}_1(u)$ as the two B -splines on two sides of \mathbf{s} .
 966 Their configuration of control points and knot vectors are the
 967 same as $\mathbf{c}(u)$'s. We denote the start and end points of \mathbf{l}_m as \mathbf{l}_m^s \mathbf{l}_m^e
 968 respectively. The fitting problem can be formulated as follows:

$$969 \arg \min_{C_0, C_1} \left(\sum_{m=0}^M \|\mathbf{c}_0(u_m) - \mathbf{l}_m^s + \mathbf{c}_1(u_m) - \mathbf{l}_m^e\|_2^2 + \beta \sum_{0 \leq i \leq j \leq n} \|\mathbf{r}_i - \mathbf{r}_j\|_2^2 \right), \quad (B.2)$$

970 where C_0 and C_1 are the control points of $\mathbf{c}_0(u)$ and $\mathbf{c}_1(u)$, \mathbf{r}_i is the
 971 difference between the corresponding control points in C_0 and C_1 .
 972 We set $\beta = 50$ in our experiments. This problem is easy to solve
 973 by using least squares.

974 **Step 3:** We perform the PCA method to generate a plane from
 975 the ruled surfaces \mathbf{s} we get in **Step 2** and then move it along its
 976 normal to find a collision-free place. To ensure the plane is long
 977 enough to cover the target patch, we set its length 1.2 times $\mathbf{c}(u)$'s.
 978 Finally, we uniformly sample along the two perpendicular edges
 979 to the ruling direction to generate several control points on each
 side, with the number of control points on each side being the
 same as $\mathbf{c}(u)$'s, thus providing a feasible initialization.